

Customers_Attraction

November 21, 2022

1 Importing Libraries and Dataset

```
[ ]: #import the libraries that will use
import pandas as pd
from google.colab import files
import io
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn import preprocessing
import numpy as np
import copy
import matplotlib.style as style
import os
import math
from scipy import stats
from collections import Counter
from sklearn.datasets import make_classification
from sklearn.model_selection import train_test_split as tts
from sklearn.decomposition import PCA
from sklearn.neighbors import KNeighborsClassifier as KNN
from sklearn.metrics import classification_report
from imblearn.over_sampling import SMOTE
from imblearn.pipeline import Pipeline
from sklearn.metrics import roc_auc_score
```

```
[ ]: pip install dython
```

Looking in indexes: <https://pypi.org/simple>, <https://us-python.pkg.dev/colab-wheels/public/simple/>

Collecting dython

Downloading dython-0.7.2-py3-none-any.whl (22 kB)

Requirement already satisfied: numpy>=1.19.5 in /usr/local/lib/python3.7/dist-packages (from dython) (1.21.6)

Requirement already satisfied: scipy>=1.7.1 in /usr/local/lib/python3.7/dist-packages (from dython) (1.7.3)

Requirement already satisfied: pandas>=1.3.2 in /usr/local/lib/python3.7/dist-packages (from dython) (1.3.5)

```

Collecting scikit-plot>=0.3.7
  Downloading scikit_plot-0.3.7-py3-none-any.whl (33 kB)
Collecting matplotlib>=3.4.3
  Downloading
matplotlib-3.5.3-cp37-cp37m-manylinux_2_5_x86_64.manylinux1_x86_64.whl (11.2 MB)
    |                                     | 11.2 MB 21.0 MB/s
Collecting psutil>=5.9.1
  Downloading psutil-5.9.4-cp36-abi3-manylinux_2_12_x86_64.manylinux2010_x86_64.
manylinux_2_17_x86_64.manylinux2014_x86_64.whl (280 kB)
    |                                     | 280 kB 13.2 MB/s
Requirement already satisfied: seaborn>=0.11.0 in
/usr/local/lib/python3.7/dist-packages (from dython) (0.11.2)
Requirement already satisfied: scikit-learn>=0.24.2 in
/usr/local/lib/python3.7/dist-packages (from dython) (1.0.2)
Requirement already satisfied: pyparsing>=2.2.1 in
/usr/local/lib/python3.7/dist-packages (from matplotlib>=3.4.3->dython) (3.0.9)
Requirement already satisfied: cyclor>=0.10 in /usr/local/lib/python3.7/dist-
packages (from matplotlib>=3.4.3->dython) (0.11.0)
Requirement already satisfied: python-dateutil>=2.7 in
/usr/local/lib/python3.7/dist-packages (from matplotlib>=3.4.3->dython) (2.8.2)
Requirement already satisfied: kiwisolver>=1.0.1 in
/usr/local/lib/python3.7/dist-packages (from matplotlib>=3.4.3->dython) (1.4.4)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.7/dist-
packages (from matplotlib>=3.4.3->dython) (21.3)
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.7/dist-
packages (from matplotlib>=3.4.3->dython) (7.1.2)
Collecting fonttools>=4.22.0
  Downloading fonttools-4.38.0-py3-none-any.whl (965 kB)
    |                                     | 965 kB 29.0 MB/s
Requirement already satisfied: typing-extensions in
/usr/local/lib/python3.7/dist-packages (from
kiwisolver>=1.0.1->matplotlib>=3.4.3->dython) (4.1.1)
Requirement already satisfied: pytz>=2017.3 in /usr/local/lib/python3.7/dist-
packages (from pandas>=1.3.2->dython) (2022.6)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.7/dist-
packages (from python-dateutil>=2.7->matplotlib>=3.4.3->dython) (1.15.0)
Requirement already satisfied: threadpoolctl>=2.0.0 in
/usr/local/lib/python3.7/dist-packages (from scikit-learn>=0.24.2->dython)
(3.1.0)
Requirement already satisfied: joblib>=0.11 in /usr/local/lib/python3.7/dist-
packages (from scikit-learn>=0.24.2->dython) (1.2.0)
Installing collected packages: fonttools, matplotlib, scikit-plot, psutil,
dython
  Attempting uninstall: matplotlib
    Found existing installation: matplotlib 3.2.2
    Uninstalling matplotlib-3.2.2:
      Successfully uninstalled matplotlib-3.2.2
  Attempting uninstall: psutil

```

```
Found existing installation: psutil 5.4.8
Uninstalling psutil-5.4.8:
  Successfully uninstalled psutil-5.4.8
Successfully installed dython-0.7.2 fonttools-4.38.0 matplotlib-3.5.3
psutil-5.9.4 scikit-plot-0.3.7
```

```
[ ]: from google.colab import drive
drive.mount('/content/drive')
```

```
[ ]: pip install matplotlib==3.1.1
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-
wheels/public/simple/
```

```
Collecting matplotlib==3.1.1
```

```
  Downloading matplotlib-3.1.1-cp37-cp37m-manylinux1_x86_64.whl (13.1 MB)
```

```
    |                                     | 13.1 MB 9.1 MB/s
```

```
Requirement already satisfied: numpy>=1.11 in
```

```
/usr/local/lib/python3.7/dist-packages (from matplotlib==3.1.1) (1.21.6)
```

```
Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in
```

```
/usr/local/lib/python3.7/dist-packages (from matplotlib==3.1.1) (3.0.9)
```

```
Requirement already satisfied: cyycler>=0.10 in /usr/local/lib/python3.7/dist-
packages (from matplotlib==3.1.1) (0.11.0)
```

```
Requirement already satisfied: python-dateutil>=2.1 in
```

```
/usr/local/lib/python3.7/dist-packages (from matplotlib==3.1.1) (2.8.2)
```

```
Requirement already satisfied: kiwisolver>=1.0.1 in
```

```
/usr/local/lib/python3.7/dist-packages (from matplotlib==3.1.1) (1.4.4)
```

```
Requirement already satisfied: typing-extensions in
```

```
/usr/local/lib/python3.7/dist-packages (from
```

```
kiwisolver>=1.0.1->matplotlib==3.1.1) (4.1.1)
```

```
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.7/dist-
packages (from python-dateutil>=2.1->matplotlib==3.1.1) (1.15.0)
```

```
Installing collected packages: matplotlib
```

```
  Attempting uninstall: matplotlib
```

```
    Found existing installation: matplotlib 3.5.3
```

```
  Uninstalling matplotlib-3.5.3:
```

```
    Successfully uninstalled matplotlib-3.5.3
```

```
ERROR: pip's dependency resolver does not currently take into account all
```

```
the packages that are installed. This behaviour is the source of the following
dependency conflicts.
```

```
dython 0.7.2 requires matplotlib>=3.4.3, but you have matplotlib 3.1.1 which is
incompatible.
```

```
Successfully installed matplotlib-3.1.1
```

```
[ ]: # Read the data with the Pandas library in this stage
```

```
data = pd.read_csv('https://raw.githubusercontent.com/AlaaAli968/Bank-Churn/
↳main/BankChurners%20(1).csv', sep = ',')
```

2 Exploratory Data Analysis

```
[ ]: # To check the datatypes as we can see do not have any null value
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10127 entries, 0 to 10126
Data columns (total 21 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   CLIENTNUM                            10127 non-null  int64
1   Attrition_Flag                       10127 non-null  object
2   Customer_Age                         10127 non-null  int64
3   Gender                               10127 non-null  object
4   Dependent_count                      10127 non-null  int64
5   Education_Level                     10127 non-null  object
6   Marital_Status                      10127 non-null  object
7   Income_Category                     10127 non-null  object
8   Card_Category                       10127 non-null  object
9   Months_on_book                      10127 non-null  int64
10  Total_Relationship_Count             10127 non-null  int64
11  Months_Inactive_12_mon               10127 non-null  int64
12  Contacts_Count_12_mon               10127 non-null  int64
13  Credit_Limit                        10127 non-null  float64
14  Total_Revolving_Bal                 10127 non-null  int64
15  Avg_Open_To_Buy                     10127 non-null  float64
16  Total_Amt_Chng_Q4_Q1                10127 non-null  float64
17  Total_Trans_Amt                     10127 non-null  int64
18  Total_Trans_Ct                      10127 non-null  int64
19  Total_Ct_Chng_Q4_Q1                 10127 non-null  float64
20  Avg_Utilization_Ratio                10127 non-null  float64
dtypes: float64(5), int64(10), object(6)
memory usage: 1.6+ MB
```

```
[ ]: # To check the data we can use the head() function to see first 5 rows.
data.head()
```

```
[ ]:  CLIENTNUM    Attrition_Flag  Customer_Age  Gender  Dependent_count  \
0   768805383  Existing Customer         45      M              3
1   818770008  Existing Customer         49      F              5
2   713982108  Existing Customer         51      M              3
3   769911858  Existing Customer         40      F              4
```

4 709106358 Existing Customer 40 M 3

| | Education_Level | Marital_Status | Income_Category | Card_Category | \ |
|---|-----------------|----------------|-----------------|---------------|---|
| 0 | High School | Married | \$60K - \$80K | Blue | |
| 1 | Graduate | Single | Less than \$40K | Blue | |
| 2 | Graduate | Married | \$80K - \$120K | Blue | |
| 3 | High School | Unknown | Less than \$40K | Blue | |
| 4 | Uneducated | Married | \$60K - \$80K | Blue | |

| | Months_on_book | ... | Months_Inactive_12_mon | Contacts_Count_12_mon | \ |
|---|----------------|-----|------------------------|-----------------------|---|
| 0 | 39 | ... | 1 | 3 | |
| 1 | 44 | ... | 1 | 2 | |
| 2 | 36 | ... | 1 | 0 | |
| 3 | 34 | ... | 4 | 1 | |
| 4 | 21 | ... | 1 | 0 | |

| | Credit_Limit | Total_Revolving_Bal | Avg_Open_To_Buy | Total_Amt_Chng_Q4_Q1 | \ |
|---|--------------|---------------------|-----------------|----------------------|---|
| 0 | 12691.0 | 777 | 11914.0 | 1.335 | |
| 1 | 8256.0 | 864 | 7392.0 | 1.541 | |
| 2 | 3418.0 | 0 | 3418.0 | 2.594 | |
| 3 | 3313.0 | 2517 | 796.0 | 1.405 | |
| 4 | 4716.0 | 0 | 4716.0 | 2.175 | |

| | Total_Trans_Amt | Total_Trans_Ct | Total_Ct_Chng_Q4_Q1 | Avg_Utilization_Ratio |
|---|-----------------|----------------|---------------------|-----------------------|
| 0 | 1144 | 42 | 1.625 | 0.061 |
| 1 | 1291 | 33 | 3.714 | 0.105 |
| 2 | 1887 | 20 | 2.333 | 0.000 |
| 3 | 1171 | 20 | 2.333 | 0.760 |
| 4 | 816 | 28 | 2.500 | 0.000 |

[5 rows x 21 columns]

```
[ ]: d=data.describe().T
d
```

```
[ ]:
count      mean      std      min  \
CLIENTNUM  10127.0  7.391776e+08  3.690378e+07  708082083.0
Customer_Age  10127.0  4.632596e+01  8.016814e+00    26.0
Dependent_count  10127.0  2.346203e+00  1.298908e+00     0.0
Months_on_book  10127.0  3.592841e+01  7.986416e+00    13.0
Total_Relationship_Count  10127.0  3.812580e+00  1.554408e+00     1.0
Months_Inactive_12_mon  10127.0  2.341167e+00  1.010622e+00     0.0
Contacts_Count_12_mon  10127.0  2.455317e+00  1.106225e+00     0.0
Credit_Limit  10127.0  8.631954e+03  9.088777e+03   1438.3
Total_Revolving_Bal  10127.0  1.162814e+03  8.149873e+02     0.0
Avg_Open_To_Buy  10127.0  7.469140e+03  9.090685e+03     3.0
Total_Amt_Chng_Q4_Q1  10127.0  7.599407e-01  2.192068e-01     0.0
```

| | | | | |
|-----------------------|---------|--------------|--------------|-------|
| Total_Trans_Amt | 10127.0 | 4.404086e+03 | 3.397129e+03 | 510.0 |
| Total_Trans_Ct | 10127.0 | 6.485869e+01 | 2.347257e+01 | 10.0 |
| Total_Ct_Chng_Q4_Q1 | 10127.0 | 7.122224e-01 | 2.380861e-01 | 0.0 |
| Avg_Utilization_Ratio | 10127.0 | 2.748936e-01 | 2.756915e-01 | 0.0 |

| | 25% | 50% | 75% | \ |
|--------------------------|--------------|--------------|--------------|---|
| CLIENTNUM | 7.130368e+08 | 7.179264e+08 | 7.731435e+08 | |
| Customer_Age | 4.100000e+01 | 4.600000e+01 | 5.200000e+01 | |
| Dependent_count | 1.000000e+00 | 2.000000e+00 | 3.000000e+00 | |
| Months_on_book | 3.100000e+01 | 3.600000e+01 | 4.000000e+01 | |
| Total_Relationship_Count | 3.000000e+00 | 4.000000e+00 | 5.000000e+00 | |
| Months_Inactive_12_mon | 2.000000e+00 | 2.000000e+00 | 3.000000e+00 | |
| Contacts_Count_12_mon | 2.000000e+00 | 2.000000e+00 | 3.000000e+00 | |
| Credit_Limit | 2.555000e+03 | 4.549000e+03 | 1.106750e+04 | |
| Total_Revolving_Bal | 3.590000e+02 | 1.276000e+03 | 1.784000e+03 | |
| Avg_Open_To_Buy | 1.324500e+03 | 3.474000e+03 | 9.859000e+03 | |
| Total_Amt_Chng_Q4_Q1 | 6.310000e-01 | 7.360000e-01 | 8.590000e-01 | |
| Total_Trans_Amt | 2.155500e+03 | 3.899000e+03 | 4.741000e+03 | |
| Total_Trans_Ct | 4.500000e+01 | 6.700000e+01 | 8.100000e+01 | |
| Total_Ct_Chng_Q4_Q1 | 5.820000e-01 | 7.020000e-01 | 8.180000e-01 | |
| Avg_Utilization_Ratio | 2.300000e-02 | 1.760000e-01 | 5.030000e-01 | |

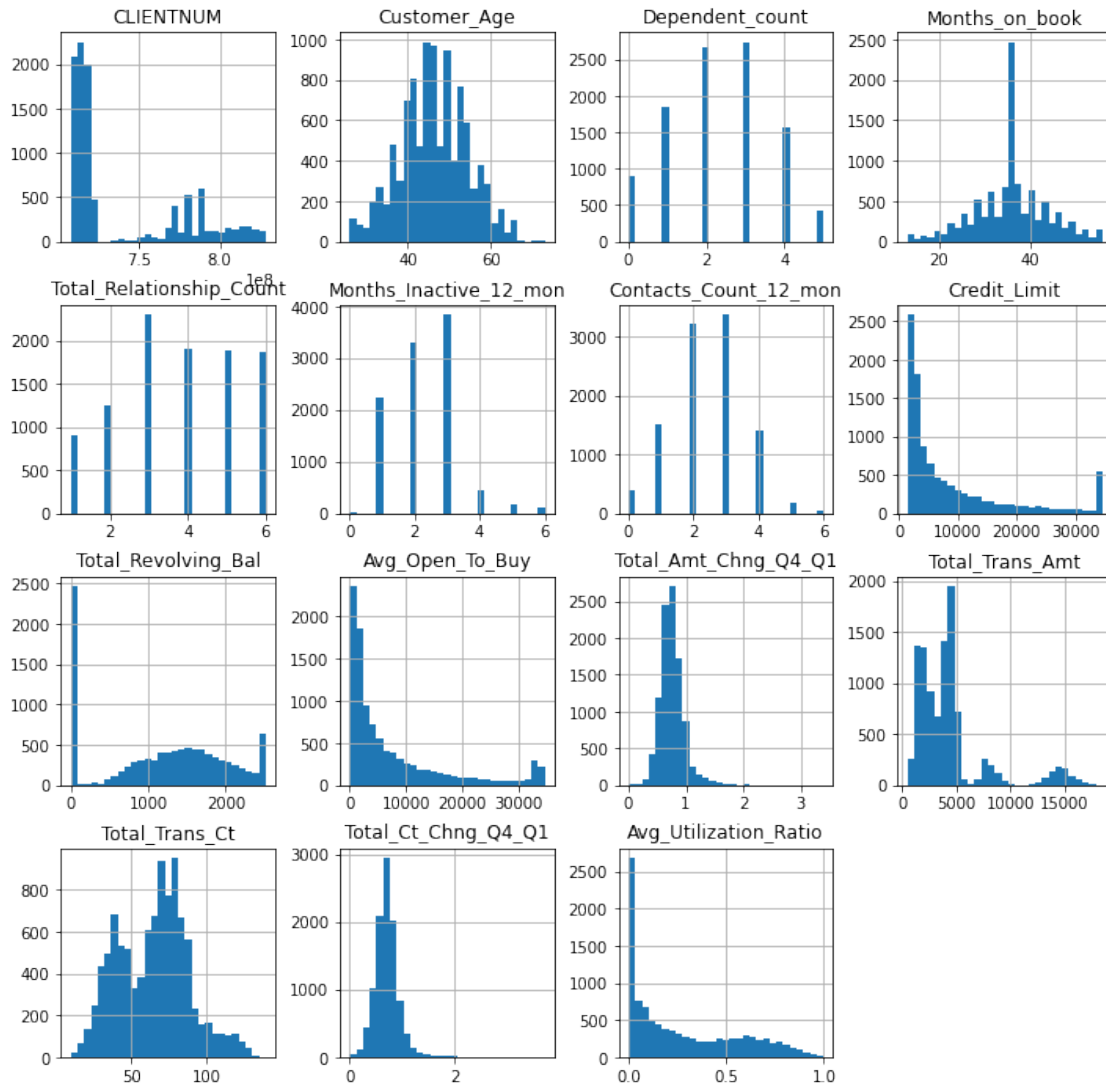
| | max |
|--------------------------|--------------|
| CLIENTNUM | 8.283431e+08 |
| Customer_Age | 7.300000e+01 |
| Dependent_count | 5.000000e+00 |
| Months_on_book | 5.600000e+01 |
| Total_Relationship_Count | 6.000000e+00 |
| Months_Inactive_12_mon | 6.000000e+00 |
| Contacts_Count_12_mon | 6.000000e+00 |
| Credit_Limit | 3.451600e+04 |
| Total_Revolving_Bal | 2.517000e+03 |
| Avg_Open_To_Buy | 3.451600e+04 |
| Total_Amt_Chng_Q4_Q1 | 3.397000e+00 |
| Total_Trans_Amt | 1.848400e+04 |
| Total_Trans_Ct | 1.390000e+02 |
| Total_Ct_Chng_Q4_Q1 | 3.714000e+00 |
| Avg_Utilization_Ratio | 9.990000e-01 |

```
[ ]: # Double check for the null-values
data.isnull().sum()
```

```
[ ]: CLIENTNUM      0
Attrition_Flag     0
Customer_Age      0
Gender            0
Dependent_count   0
```

| | |
|--------------------------|---|
| Education_Level | 0 |
| Marital_Status | 0 |
| Income_Category | 0 |
| Card_Category | 0 |
| Months_on_book | 0 |
| Total_Relationship_Count | 0 |
| Months_Inactive_12_mon | 0 |
| Contacts_Count_12_mon | 0 |
| Credit_Limit | 0 |
| Total_Revolving_Bal | 0 |
| Avg_Open_To_Buy | 0 |
| Total_Amt_Chng_Q4_Q1 | 0 |
| Total_Trans_Amt | 0 |
| Total_Trans_Ct | 0 |
| Total_Ct_Chng_Q4_Q1 | 0 |
| Avg_Utilization_Ratio | 0 |
| dtype: int64 | |

```
[ ]: # distribution of numerical features
axList = data.hist(bins=29, figsize = (12, 12))
plt.savefig("Hist.png")
```




```
[ ]: table=data['Attrition_Flag'].value_counts(normalize=True) * 100
print(table)
churn=data['Attrition_Flag'].value_counts()
churn
plt.figure(figsize = (6,6))
piechart=plt.pie(x=churn,labels=churn.keys(),autopct="%.1f%%")
plt.title('Proportion of Existing and Attrited Customer', fontsize = 16)
# as we see data is imbalanced so we will apply some and random oversampling
↳ techniques later to balance it before running the models
```

```
Existing Customer      83.934038
Attrited Customer      16.065962
Name: Attrition_Flag, dtype: float64
```



```
[ ]: Text(0.5, 1.0, 'Proportion of Existing and Attrited Customer')
```

Proportion of Existing and Attrited Customer



```
[ ]: new_data= copy.copy(data)

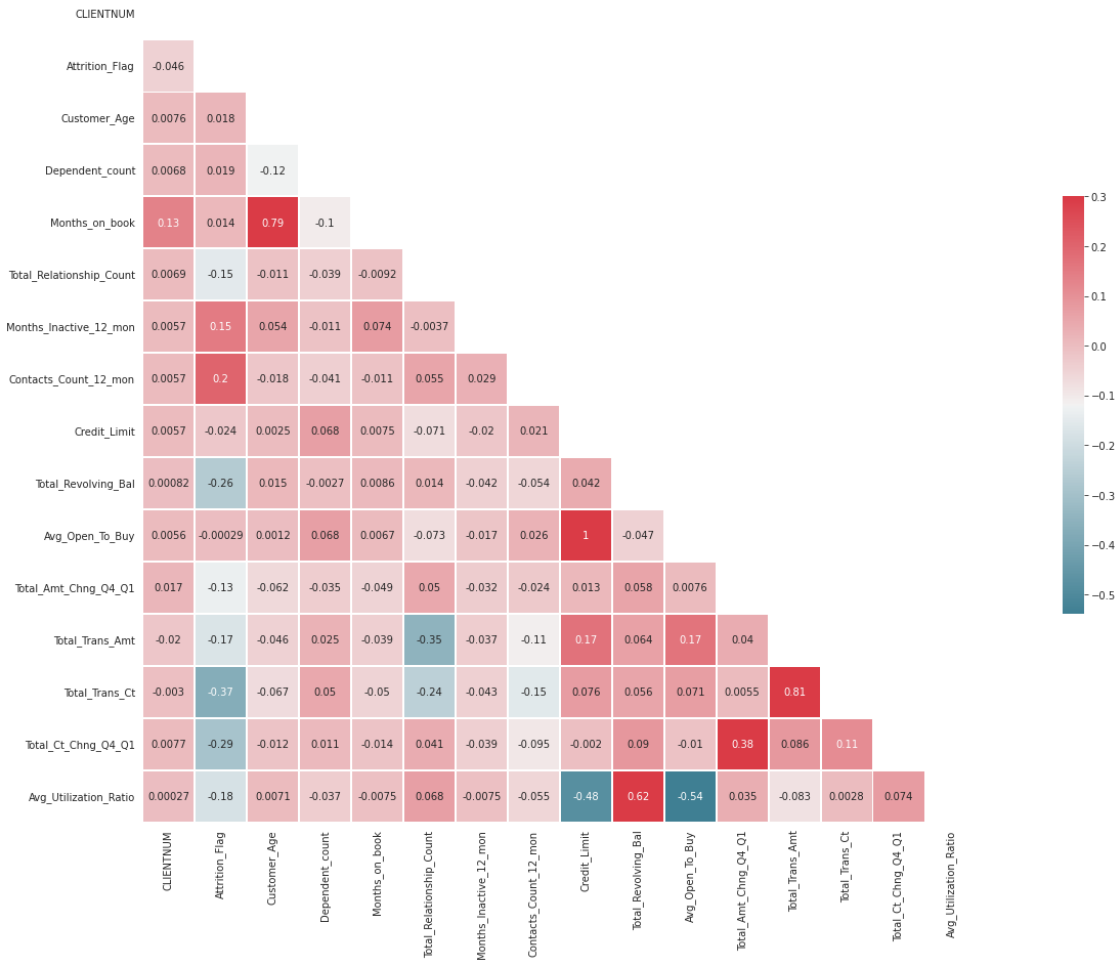
new_data['Attrition_Flag'].replace({'Existing Customer':0, 'Attrited Customer':
↪1}, inplace=True)
corr = new_data.corr()
corr

mask = np.zeros_like(corr)
mask[np.triu_indices_from(mask)] = True

# Generate a custom diverging colormap
cmap = sns.diverging_palette(220, 10, as_cmap=True)


with sns.axes_style("white"):
    # Set up the matplotlib figure
    f, ax = plt.subplots(figsize=(30, 15))
```

```
ax = sns.heatmap(corr, cmap=cmap, mask=mask, vmax=.3, square=True,
↳linewidths=.9, cbar_kws={"shrink": .5}, annot=True)
```




```
[ ]: # Customers' Distribution based on Education Level
plt.figure(figsize = (6,6))
plt.pie(data['Income_Category'].value_counts(),
        labels = ['Less than $40K', '$40K - $60K', '$80K - $120K', '$60K -
↳$80K', 'Unknown', '$120K +'],
        autopct='%1.1f%%', startangle = 90)
plt.title('Proportion of Income_Category', fontsize = 16)
plt.show()
```

Proportion of Income_Category




```
[ ]: # Customers' Distribution based on Education Level
plt.figure(figsize = (6,6))
plt.pie(data['Marital_Status'].value_counts(),
        labels = ['Married', 'Single', 'Unknown', 'Divorced'],
        autopct='%1.1f%%', startangle = 90)
plt.title('Proportion of Marital_Status', fontsize = 16)
plt.show()
```


Proportion of Marital_Status




```
[ ]: #Income_Category and Credit_Limit
sns.set_style("whitegrid")
s=sns.boxplot(x = 'Income_Category', y = 'Credit_Limit', data =_
↪data,palette="colorblind")
s.tick_params(labelsize=7.5)
```



```
[ ]: plt.scatter(data['Total_Revolving_Bal'][(data.Attrition_Flag == 'Existing_
↳Customer') | (data.Attrition_Flag == 'Attrited Customer')],
            data['Months_on_book'][(data.Attrition_Flag == 'Existing Customer')_
↳| (data.Attrition_Flag == 'Attrited Customer')],
            marker='D',
            color='lightgray',
            label='Existing Customer')
plt.scatter(data['Total_Revolving_Bal'][data.Attrition_Flag == 'Attrited_
↳Customer'],
            data['Months_on_book'][data.Attrition_Flag == 'Attrited Customer'],
            marker='+',
            color='red',
            label='Attrited Customer')
plt.xlabel('Total_Revolving_Bal')
plt.ylabel('Months_on_book')
plt.legend()
plt.show()
```



```
[ ]: #Existing Customer Attrited Customer
plt.scatter(data['Total_Trans_Ct'][(data.Attrition_Flag == 'Existing Customer') |
    ↳ (data.Attrition_Flag == 'Attrited Customer')],
            data['Credit_Limit'][(data.Attrition_Flag == 'Existing Customer') |
    ↳ (data.Attrition_Flag == 'Attrited Customer')],
            marker='D',
            color='lightgray',
            label='Existing Customer')
plt.scatter(data['Total_Trans_Ct'][data.Attrition_Flag == 'Attrited Customer'],
            data['Credit_Limit'][data.Attrition_Flag == 'Attrited Customer'],
            marker='+',
            color='red',
            label='Attrited Customer')
plt.xlabel('Total_Trans_Ct')
plt.ylabel('Credit_Limit')
plt.legend()
plt.show()
```



3 Data Preprocessing

```
[ ]: df = data.copy()
```

3.1 Converting Categorical Columns to Numeric Columns

Machine learning algorithms work best with numerical data. However, in my dataset, I have some categorical columns, I need to convert them to numeric columns.

```
[ ]: # Encoding our target: Attrition_Flag
df = df.replace({'Attrition_Flag':{'Existing Customer':0, 'Attrited Customer':
    ↳1}})
# Selecting all the categorical features to one hot encode
df_cat = df.select_dtypes('object')

# One hot encoding
df = pd.get_dummies(df, df_cat.columns, drop_first = False)
df.dtypes
```

```
[ ]: CLIENTNUM          int64
Attrition_Flag         int64
Customer_Age          int64
```

| | |
|---------------------------------|---------|
| Dependent_count | int64 |
| Months_on_book | int64 |
| Total_Relationship_Count | int64 |
| Months_Inactive_12_mon | int64 |
| Contacts_Count_12_mon | int64 |
| Credit_Limit | float64 |
| Total_Revolving_Bal | int64 |
| Avg_Open_To_Buy | float64 |
| Total_Amt_Chng_Q4_Q1 | float64 |
| Total_Trans_Amt | int64 |
| Total_Trans_Ct | int64 |
| Total_Ct_Chng_Q4_Q1 | float64 |
| Avg_Utilization_Ratio | float64 |
| Gender_F | uint8 |
| Gender_M | uint8 |
| Education_Level_College | uint8 |
| Education_Level_Doctorate | uint8 |
| Education_Level_Graduate | uint8 |
| Education_Level_High School | uint8 |
| Education_Level_Post-Graduate | uint8 |
| Education_Level_Uneducated | uint8 |
| Education_Level_Unknown | uint8 |
| Marital_Status_Divorced | uint8 |
| Marital_Status_Married | uint8 |
| Marital_Status_Single | uint8 |
| Marital_Status_Unknown | uint8 |
| Income_Category_\$120K + | uint8 |
| Income_Category_\$40K - \$60K | uint8 |
| Income_Category_\$60K - \$80K | uint8 |
| Income_Category_\$80K - \$120K | uint8 |
| Income_Category_Less than \$40K | uint8 |
| Income_Category_Unknown | uint8 |
| Card_Category_Blue | uint8 |
| Card_Category_Gold | uint8 |
| Card_Category_Platinum | uint8 |
| Card_Category_Silver | uint8 |
| dtype: object | |

4 Split Data

I split the model into 70% training and 30% testing as it's the most popular ratio

```
[ ]: from sklearn.model_selection import train_test_split
x = df.drop(['CLIENTNUM', 'Attrition_Flag'], axis = 1)
y = df['Attrition_Flag']
```



```
x_train, x_test, y_train, y_test = train_test_split(
    x,
    y,
    test_size = 0.3,
    random_state = 1999)
```

5 Model Training

After splitting the data, I will create various machine learning classifiers and identify the best model out of three.

I will train different classifiers and try to best model. We will utilize:

1. Random Forest
2. Decision Tree
3. Logistic Regression

I chose these models because they are popular in binary classification problem also as I saw in Kaggle many people who have the same dataset and used these models got more accurate results. As well as I mentioned in the literature review in the related studies section most of them used Random Forest and was the best model, those studies have binary classification problems like my problem. so based on these I think Random Forest will be the best model.

I will train these models on imbalanced data then with oversampling, I used 2 techniques for oversampling (SMOTE and Random Oversampling) and compare between them.

6 Imbalanced

6.1 Build and Evaluate the Models for Imbalanced Data

```
[ ]: from sklearn.metrics import roc_curve, \
    ↳roc_auc_score, accuracy_score, recall_score, f1_score, brier_score_loss, matthews_corrcoef

from sklearn import tree, ensemble, linear_model
#Build the 3 models
models = [tree.DecisionTreeClassifier(random_state=2002),
    ↳ensemble.RandomForestClassifier(random_state=2002),
    ↳linear_model.LogisticRegression(solver="liblinear", random_state=2002)] #
    ↳liblinear is a good choice for small dataset

name = []

Accuracy = [] #= TP+TN/TP+FP+FN+TN
```

```

Specificity=[]#true negative rate

Sensetivity=[]#true positive rate

auc=[]#evaluates the the model's performance across different threshold
f1=[]#
mcc=[]#Matthews correlation coefficient -1 to 1
brier=[]#1-Acc

for i in models:

    name.append(i.__class__.__name__)

    i.fit(x_train, y_train)

    y_predicted=i.predict(x_test)

    Accuracy.append(accuracy_score(y_test,y_predicted))

    Specificity.append(recall_score(y_test, y_predicted, pos_label=0))#true
    ↳positive rate.

    Sensetivity.append(recall_score(y_test, y_predicted, pos_label=1))#true
    ↳positive rate.

    i_probs= i.predict_proba(x_test)

    i_probs=i_probs[:,1]
    mcc.append(matthews_corrcoef(y_test,y_predicted))
    auc.append(roc_auc_score(y_test,i_probs))
    f1.append(f1_score(y_test,y_predicted))
    brier.append(brier_score_loss(y_test,y_predicted))#1-ACC

models_evaluation = pd.DataFrame({"Model": name, "MCC":mcc,"Brier":
    ↳brier,"Accuracy": Accuracy,"Sensetivity":Sensetivity,"Specificity":
    ↳Specificity,"F1":f1,"AUC":auc})

```

```
display(models_evaluation)
```

| | Model | MCC | Brier | Accuracy | Sensetivity \ |
|---|------------------------|----------|----------|----------|---------------|
| 0 | DecisionTreeClassifier | 0.742271 | 0.067456 | 0.932544 | 0.776371 |
| 1 | RandomForestClassifier | 0.830211 | 0.042777 | 0.957223 | 0.774262 |
| 2 | LogisticRegression | 0.609045 | 0.093452 | 0.906548 | 0.554852 |

| | Specificity | F1 | AUC |
|---|-------------|----------|----------|
| 0 | 0.961404 | 0.782147 | 0.868887 |
| 1 | 0.991033 | 0.849537 | 0.987695 |
| 2 | 0.971540 | 0.649383 | 0.922995 |

from the table we can see that Random Forest model the best model with really high values for Accuracy, MCC, and AUC. Also lowest value with Brier

6.2 ROC curve for Imbalanced data

Next step will show ROC curve (receiver operating characteristic curve) a graph showing the performance of a classification model at all classification thresholds(default 0.5). This curve plots two parameters:

1. True Positive Rat
2. False Positive Rate

```
[ ]: from sklearn import metrics
from sklearn import datasets
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
plt.figure(0).clf()
#fit logistic regression model and plot ROC curve
model = linear_model.LogisticRegression(solver="liblinear",random_state=2002)
model.fit(x_train, y_train)
y_pred = model.predict_proba(x_test)[:, 1]# prob number of classes
fpr, tpr, _ = metrics.roc_curve(y_test, y_pred)#fpr:False Positive Rates, tpr:
↪True Positive Rates, _:threshold
auc = round(metrics.roc_auc_score(y_test, y_pred), 4)
plt.plot(fpr,tpr,label="Imbalanced Logistic Regression, AUC="+str(auc))

#fit Decision Tree model and plot ROC curve
model = tree.DecisionTreeClassifier(random_state=2002)
model.fit(x_train, y_train)
y_pred = model.predict_proba(x_test)[:, 1]
fpr, tpr, _ = metrics.roc_curve(y_test, y_pred)
auc = round(metrics.roc_auc_score(y_test, y_pred), 4)
plt.plot(fpr,tpr,label="Imbalanced Decision Tree, AUC="+str(auc))

#fit Random Forest model and plot ROC curve
```

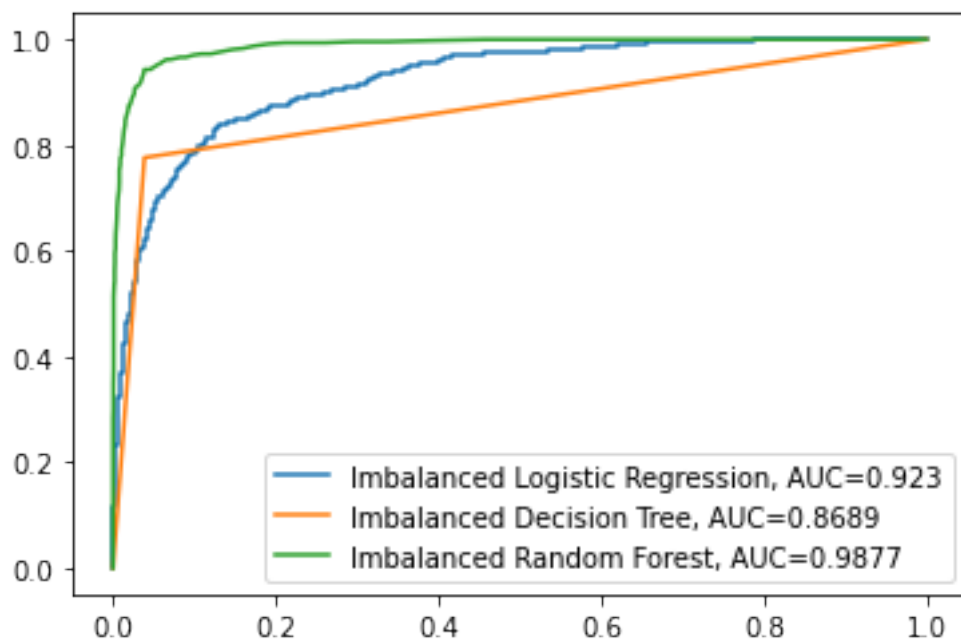
```

model = ensemble.RandomForestClassifier(random_state=2002)
model.fit(x_train, y_train)
y_pred = model.predict_proba(x_test)[: , 1]
fpr, tpr, _ = metrics.roc_curve(y_test, y_pred)
auc = round(metrics.roc_auc_score(y_test, y_pred), 4)
plt.plot(fpr,tpr,label="Imbalanced Random Forest, AUC="+str(auc))

#add legend
plt.legend()

```

```
[ ]: <matplotlib.legend.Legend at 0x7fecda00cdd0>
```



6.3 k-fold cross-validation for Imbalanced data

To check overfitting and underfitting I used Brier score which is similar to mean squared error

```

[ ]: # Evaluate a logistic regression model using repeated k-fold cross-validation

from numpy import mean

from numpy import std

from sklearn.datasets import make_classification

```

```

from sklearn.model_selection import RepeatedKFold

from sklearn.model_selection import cross_val_score

from sklearn.linear_model import LogisticRegression

# prepare the cross-validation procedure Repeats 10-Fold 3 times

cv = RepeatedKFold(n_splits=10, n_repeats=3, random_state=1)

# create models

model1 = linear_model.LogisticRegression(solver="liblinear",random_state=2002)
model2 =ensemble.RandomForestClassifier(random_state=2002)
model3 = tree.DecisionTreeClassifier(random_state=2002)

# evaluate models by Brier score which is = 1-Accuracy
scores1 = 1-cross_val_score(model1, x_train, y_train, scoring='accuracy',
    ↳cv=cv, n_jobs=-1) #n_job=-1 means using all processors
scores2 =1- cross_val_score(model2, x_train, y_train, scoring='accuracy',
    ↳cv=cv, n_jobs=-1)
scores3 = 1-cross_val_score(model3,x_train, y_train, scoring='accuracy', cv=cv,
    ↳n_jobs=-1)

#to compare with test set
log_test = 1-cross_val_score(model1, x_test, y_test, scoring='accuracy', cv=cv,
    ↳n_jobs=-1)
rf_test =1- cross_val_score(model2, x_test, y_test, scoring='accuracy', cv=cv,
    ↳n_jobs=-1)
dt_test = 1-cross_val_score(model3, x_test, y_test, scoring='accuracy', cv=cv,
    ↳n_jobs=-1)

# report performance

print('Brier for Logistic Regression : %.3f (%.3f)' %
    ↳(mean(scores1),std(scores2)))
print('Brier Random Forest: %.3f (%.3f)' % (mean(scores2), std(scores2)))
print('Brier Decision Tree: %.3f (%.3f)' % (mean(scores3), std(scores3)))
#plot performance for Logistic Regression Model
plt.plot(scores1, label='Train')
plt.plot(log_test ,label='Test')
plt.legend(loc='best')
plt.title('Logistic Regression Model')

```

```


plt.xlabel('number of Folds')
plt.ylabel('Brier Scores')
plt.savefig('number-of-cv.png')
plt.show()


#plot performance for Random Forest Model
plt.plot(scores2, label='Train')
plt.plot(rf_test ,label='Test')
plt.legend(loc='best')
plt.title('Random Forest Model')
plt.xlabel('number of Folds')
plt.ylabel('Brier Scores')
plt.savefig('number-of-cv.png')
plt.show()

#plot performance for Decision Tree Model
plt.plot(scores3, label='Train')
plt.plot(dt_test ,label='Test')
plt.legend(loc='best')
plt.title('Decision Tree Model')
plt.xlabel('number of Folds')
plt.ylabel('Brier Scores')
plt.savefig('number-of-cv.png')
plt.show()

```

Brier for Logistic Regression : 0.110 (0.007)
 Brier Random Forest: 0.046 (0.007)
 Brier Decision Tree: 0.064 (0.008)





7 SMOTE

Now I will apply SMOTE oversampling technique to fix the imbalanced data

```
[ ]: sm = SMOTE(sampling_strategy='minority', random_state=42)
x_train_res, y_train_res = sm.fit_resample(x_train, y_train)
```

```
[ ]: #Before SMOTE
unique, count = np.unique(y_train, return_counts=True)
y_train_dict_value_count = {k: v for (k, v) in zip(unique, count)}
print(y_train_dict_value_count)

#After SMOTE
unique, count = np.unique(y_train_res, return_counts=True)
y_train_dict_value_count = {k: v for (k, v) in zip(unique, count)}
print(y_train_dict_value_count)
```

```
{0: 5935, 1: 1153}
{0: 5935, 1: 5935}
```


7.1 Build and Evaluate the Models for balanced Data (SMOTE)

```
[ ]: from sklearn.metrics import roc_curve, \
      →roc_auc_score, accuracy_score, recall_score, f1_score, matthews_corrcoef

from sklearn import tree, ensemble, linear_model

models = [tree.DecisionTreeClassifier(random_state=2002),
          ensemble.RandomForestClassifier(random_state=2002),
          linear_model.LogisticRegression(solver="liblinear", random_state=2002)]

name = []

Accuracy = []

Specificity = []

Sensitivity = []

auc = []

f1 = []

mcc = []
brier = []

for i in models:

    name.append(i.__class__.__name__)

    i.fit(x_train_res, y_train_res)

    y_predicted = i.predict(x_test)

    matthews_corrcoef(y_test, y_predicted)

    Accuracy.append(accuracy_score(y_test, y_predicted))

    Specificity.append(recall_score(y_test, y_predicted, pos_label=0))

    Sensitivity.append(recall_score(y_test, y_predicted, pos_label=1))
    mcc.append(matthews_corrcoef(y_test, y_predicted))

    i_probs = i.predict_proba(x_test)

    i_probs = i_probs[:, 1]
```

```

auc.append(roc_auc_score(y_test,i_probs))
f1.append(f1_score(y_test,y_predicted))
brier.append(brier_score_loss(y_test,y_predicted))

models_evaluation = pd.DataFrame({"Model": name,"MCC":mcc ,"Brier":
    ↳brier,"Accuracy": Accuracy,"Sensetivity":Sensetivity,"Specificity":
    ↳Specificity,"F1":f1,"AUC":auc}).style.set_caption("After SMOTE OverSampling")

display(models_evaluation)

```

```
<pandas.io.formats.style.Styler at 0x7fec91a4890>
```

As we can see, Random Forest is still the best model, but there is a slight decrease in the Accuracy, Auc and true negative rate after using SMOTE. On the other hand F1, true positive rate and MCC increased.

7.2 ROC curve for balanced Data (SMOTE)

```

[ ]: from sklearn import metrics
from sklearn import datasets
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt

#fit logistic regression model and plot ROC curve
model = linear_model.LogisticRegression(solver="liblinear",random_state=2002)
model.fit(x_train_res, y_train_res)
y_pred = model.predict_proba(x_test)[:, 1]
fpr, tpr, _ = metrics.roc_curve(y_test, y_pred)
auc = round(metrics.roc_auc_score(y_test, y_pred), 4)
plt.plot(fpr,tpr,label="SMOTE Logistic Regression, AUC="+str(auc))

#fit Decision Tree model and plot ROC curve
model = tree.DecisionTreeClassifier(random_state=2002)
model.fit(x_train_res, y_train_res)
y_pred = model.predict_proba(x_test)[:, 1]
fpr, tpr, _ = metrics.roc_curve(y_test, y_pred)
auc = round(metrics.roc_auc_score(y_test, y_pred), 4)
plt.plot(fpr,tpr,label="SMOTE Decision Tree, AUC="+str(auc))

#fit Random Forest model and plot ROC curve

```

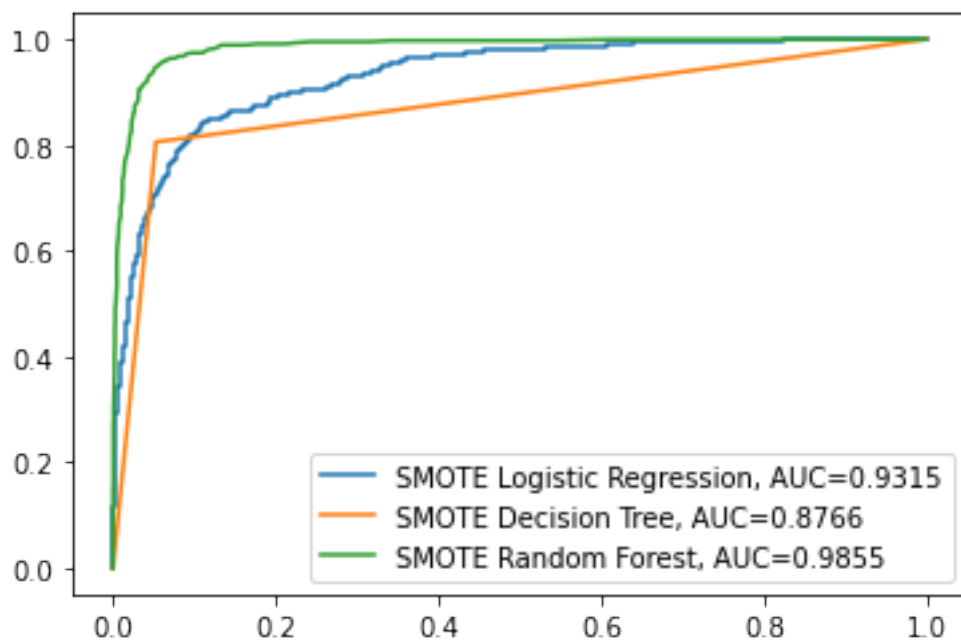
```

model = ensemble.RandomForestClassifier(random_state=2002)
model.fit(x_train_res, y_train_res)
y_pred = model.predict_proba(x_test)[: , 1]
fpr, tpr, _ = metrics.roc_curve(y_test, y_pred) #thrushhold
auc = round(metrics.roc_auc_score(y_test, y_pred), 4)
plt.plot(fpr,tpr,label="SMOTE Random Forest, AUC="+str(auc))

#add legend
plt.legend()

```

```
[ ]: <matplotlib.legend.Legend at 0x7fec85f0190>
```



7.3 k-fold cross-validation for balanced data(SMOTE)

```

[ ]: from sklearn.datasets import make_classification

from sklearn.model_selection import RepeatedKFold

from sklearn.model_selection import cross_val_score

from sklearn.linear_model import LogisticRegression

```

```

# prepare the cross-validation procedure Repeats 10-Fold 3 times

cv = RepeatedKfold(n_splits=10, n_repeats=3, random_state=1)#

# create models

model1 = linear_model.LogisticRegression(solver="liblinear",random_state=2002)
model2 =ensemble.RandomForestClassifier(random_state=2002)
model3 = tree.DecisionTreeClassifier(random_state=2002)

# evaluate models by Brier score which is = 1-Accuracy
scores1 = 1-cross_val_score(model1, x_train_res, y_train_res,␣
    ↳scoring='accuracy', cv=cv, n_jobs=-1) #n_job=-1 -1 means using all processors
scores2 =1- cross_val_score(model2, x_train_res, y_train_res,␣
    ↳scoring='accuracy', cv=cv, n_jobs=-1)
scores3 = 1-cross_val_score(model3,x_train_res, y_train_res,␣
    ↳scoring='accuracy', cv=cv, n_jobs=-1)
# to compare with test set
log_test = 1-cross_val_score(model1, x_test, y_test, scoring='accuracy', cv=cv,␣
    ↳n_jobs=-1)
rf_test =1- cross_val_score(model2, x_test, y_test, scoring='accuracy', cv=cv,␣
    ↳n_jobs=-1)
dt_test = 1-cross_val_score(model3, x_test, y_test, scoring='accuracy', cv=cv,␣
    ↳n_jobs=-1)

#print(scores)

# report performance

print('Brier for Logistic Regression : %.3f (%.3f)' %␣
    ↳(mean(scores1),std(scores2)))
print('Brier for Random Forest: %.3f (%.3f)' % (mean(scores2), std(scores2)))
print('Brier for Decision Tree: %.3f (%.3f)' % (mean(scores3), std(scores3)))
#plot performance for Logistic Regression Model
plt.plot(scores1, label='Train')
plt.plot(log_test ,label='Test')
plt.legend(loc='best')
plt.title('Logistic Regression Model')
plt.xlabel('number of Folds')
plt.ylabel('Brier Scores')
plt.savefig('number-of-cv.png')
plt.show()
#plot performance for Random Forest Model

plt.plot(scores2, label='Train')

```

```


plt.plot(rf_test ,label='Test')
plt.legend(loc='best')
plt.title('Random Forest Model')
plt.xlabel('number of Folds')
plt.ylabel('Brier Scores')
plt.savefig('number-of-cv.png')
plt.show()
#plot performance for Decision Tree Model
plt.plot(scores3, label='Train')
plt.plot(dt_test ,label='Test')
plt.legend(loc='best')
plt.title('Decision Tree Model')
plt.xlabel('number of Folds')
plt.ylabel('Brier Scores')
plt.savefig('number-of-cv.png')
plt.show()

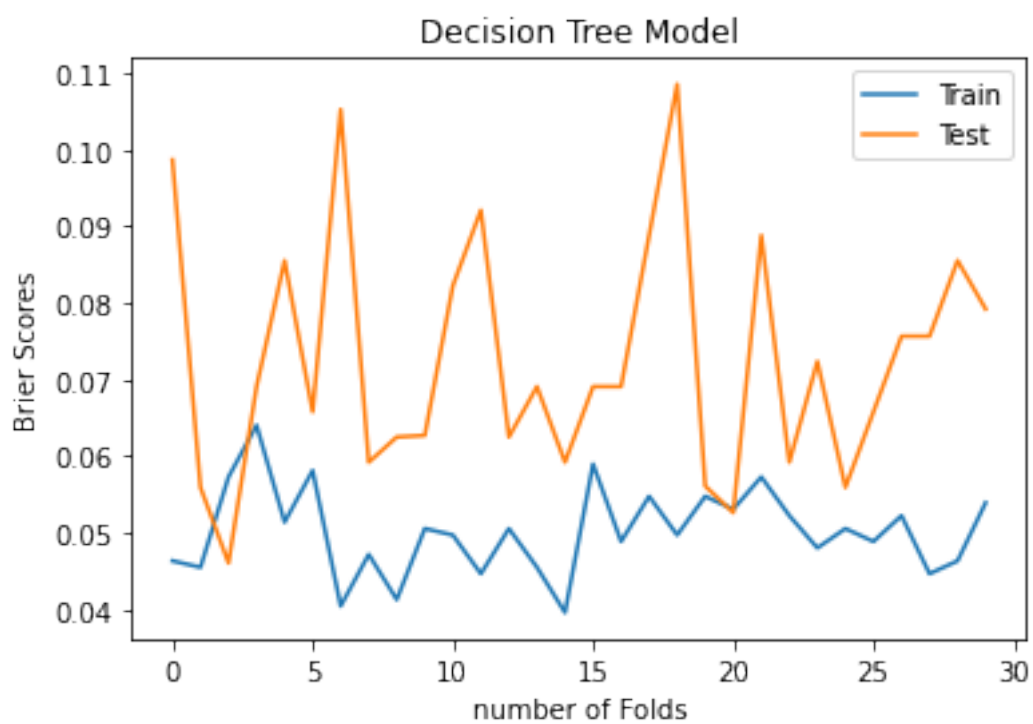
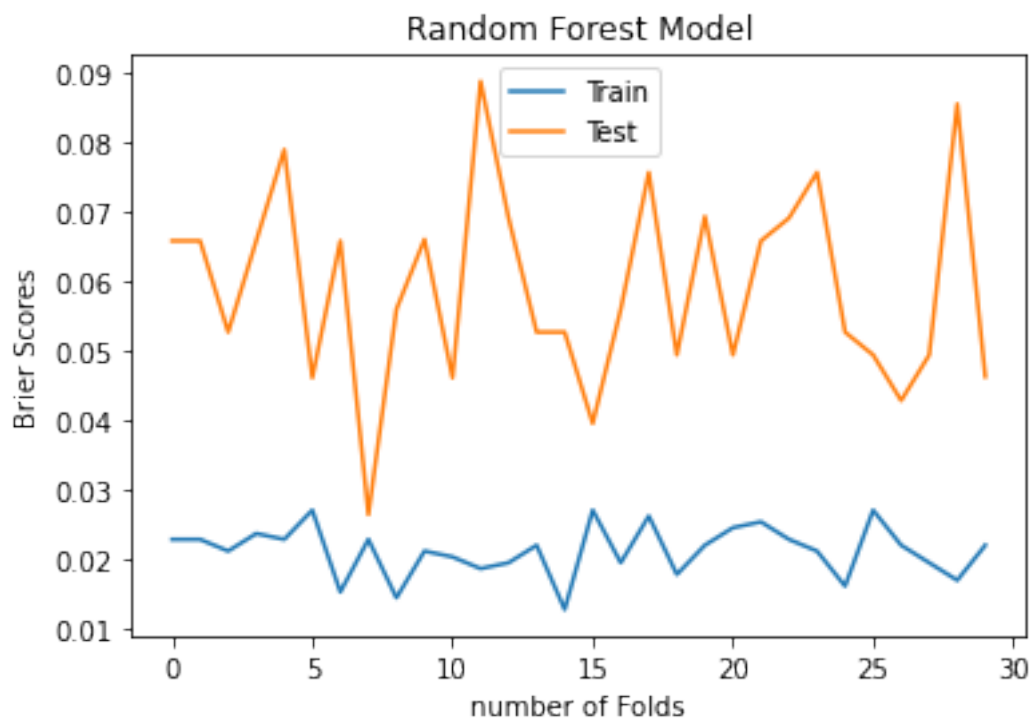
```

Brier for Logistic Regression : 0.071 (0.004)

Brier for Random Forest: 0.021 (0.004)

Brier for Decision Tree: 0.050 (0.006)





8 Random OverSampling

```
[ ]: from sklearn.datasets import make_classification
      from imblearn.over_sampling import RandomOverSampler

      # define oversampling strategy
      oversample = RandomOverSampler(sampling_strategy='minority',random_state=2002)
      # fit and apply the transform
      x_over, y_over = oversample.fit_resample(x_train, y_train)
```

```
[ ]: x_over.shape[1]
```

```
[ ]: 37
```

```
[ ]: #Before Random OverSampling
      unique,count=np.unique(y_train,return_counts=True)
      y_train_dict_value_count={k:v for (k,v) in zip(unique,count)}
      print(y_train_dict_value_count)

      #After Random OverSampling
      unique,count=np.unique(y_over,return_counts=True)
      y_train_dict_value_count={k:v for (k,v) in zip(unique,count)}
      print(y_train_dict_value_count)
```

```
{0: 5935, 1: 1153}
```

```
{0: 5935, 1: 5935}
```

```
[ ]: from sklearn.metrics import roc_curve, \
      ↪roc_auc_score,accuracy_score,recall_score,f1_score

      from sklearn import tree,ensemble

      models = [tree.DecisionTreeClassifier(random_state=2002),
                  ensemble.RandomForestClassifier(random_state=2002),
                  linear_model.LogisticRegression(solver="liblinear",random_state=2002)]

      name = []

      Accuracy = []

      Specificity=[]

      Sensetivity=[]

      auc=[]

      f1=[]
```

```

mcc=[]

brier=[]

for i in models:

    name.append(i.__class__.__name__)

    i.fit(x_over, y_over)

    y_predicted=i.predict(x_test)

    Accuracy.append(accuracy_score(y_test,y_predicted))

    Specificity.append(recall_score(y_test, y_predicted, pos_label=0))

    Sensetivity.append(recall_score(y_test, y_predicted, pos_label=1))

    i_probs= i.predict_proba(x_test)

    i_probs=i_probs[:,1]
    mcc.append(matthews_corrcoef(y_test,y_predicted))

    auc.append(roc_auc_score(y_test,i_probs))
    f1.append(f1_score(y_test,y_predicted))
    brier.append(brier_score_loss(y_test,y_predicted))

models_evaluation = pd.DataFrame({"Model": name,"MCC":mcc,"Brier":brier,
    ↳ "Accuracy": Accuracy,"Sensetivity":Sensetivity,"Specificity":
    ↳ Specificity,"F1":f1,"AUC":auc}).style.set_caption("After Random
    ↳ OverSampling")

display(models_evaluation)

```

```
<pandas.io.formats.style.Styler at 0x7fec91ef110>
```


Random Forest with Random Oversampling gave the best results compared with SMOTE and imbalanced data with all results. MCC 0.856629 which is close to 1 our target. Brier scores very close to zero. Accuracy, Sensitivity, Specificity, F1 and AUC increased.

8.1 ROC curve for balanced Data (Random OverSampling)

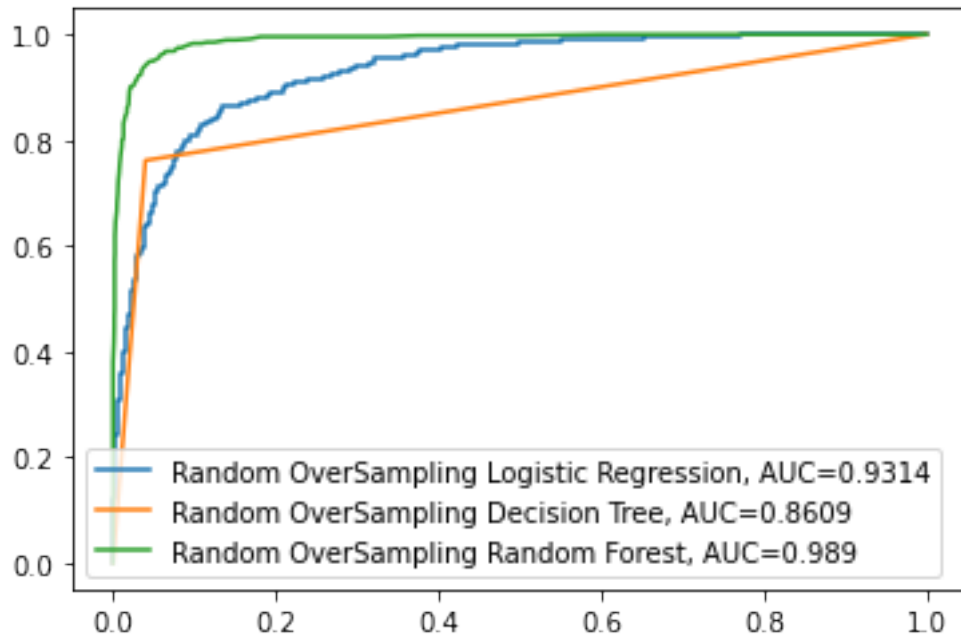
```
[ ]: #plt.figure(0).clf()
#fit logistic regression model and plot ROC curve
model = linear_model.LogisticRegression(solver="liblinear",random_state=2002)
model.fit(x_over, y_over)
y_pred = model.predict_proba(x_test)[:, 1]
fpr, tpr, _ = metrics.roc_curve(y_test, y_pred)
auc = round(metrics.roc_auc_score(y_test, y_pred), 4)
plt.plot(fpr,tpr,label="Random OverSampling Logistic Regression, AUC="+str(auc))

#fit Decision Tree model and plot ROC curve
model = tree.DecisionTreeClassifier(random_state=2002)
model.fit(x_over, y_over)
y_pred = model.predict_proba(x_test)[:, 1]
fpr, tpr, _ = metrics.roc_curve(y_test, y_pred)
auc = round(metrics.roc_auc_score(y_test, y_pred), 4)
plt.plot(fpr,tpr,label="Random OverSampling Decision Tree, AUC="+str(auc))

#fit Random Forest model and plot ROC curve
model = ensemble.RandomForestClassifier(random_state=2002)
model.fit(x_over, y_over)
y_pred = model.predict_proba(x_test)[:, 1]
fpr, tpr, _ = metrics.roc_curve(y_test, y_pred)
auc = round(metrics.roc_auc_score(y_test, y_pred), 4)
plt.plot(fpr,tpr,label="Random OverSampling Random Forest, AUC="+str(auc))

#add legend
plt.legend()
```

```
[ ]: <matplotlib.legend.Legend at 0x7fec8603a50>
```



```
[ ]: # evaluate a logistic regression model using repeated k-fold cross-validation

from numpy import mean

from numpy import std

from sklearn.datasets import make_classification

from sklearn.model_selection import RepeatedKFold

from sklearn.model_selection import cross_val_score

from sklearn.linear_model import LogisticRegression

# prepare the cross-validation procedure

cv = RepeatedKFold(n_splits=10, n_repeats=3, random_state=1)

# create models
model1 = linear_model.LogisticRegression(solver="liblinear", random_state=2002)
model2 = ensemble.RandomForestClassifier(random_state=2002)
model3 = tree.DecisionTreeClassifier(random_state=2002)

# evaluate models by Brier score which is = 1-Accuracy
```

```

scores1 = 1-cross_val_score(model1, x_over, y_over, scoring='accuracy', cv=cv,
    ↪n_jobs=-1)
scores2 = 1-cross_val_score(model2, x_over, y_over, scoring='accuracy', cv=cv,
    ↪n_jobs=-1)
scores3 = 1-cross_val_score(model3, x_over, y_over, scoring='accuracy', cv=cv,
    ↪n_jobs=-1)

log_test = 1-cross_val_score(model1, x_test, y_test, scoring='accuracy', cv=cv,
    ↪n_jobs=-1)
rf_test = 1-cross_val_score(model2, x_test, y_test, scoring='accuracy', cv=cv,
    ↪n_jobs=-1)
dt_test = 1-cross_val_score(model3, x_test, y_test, scoring='accuracy', cv=cv,
    ↪n_jobs=-1)

# print(scores)

# report performance

print('Average Brier for Logistic Regression : %.3f (%.3f)' %
    ↪(mean(scores1),std(scores2)))
print('Average Brier Random Forest: %.3f (%.3f)' % (mean(scores2),
    ↪std(scores2)))
print('Average Brier Decision Tree: %.3f (%.3f)' % (mean(scores3),
    ↪std(scores3)))

plt.plot(scores1, label='over')
plt.plot(log_test, label='test')
plt.legend(loc='best')
plt.title('Logistic Regression')
plt.xlabel('number of Folds')
plt.ylabel('mean classification score')
plt.savefig('number-of-cv.png')
plt.show()

plt.plot(scores2, label='over')
plt.plot(rf_test, label='test')
plt.legend(loc='best')
plt.title('RF')
plt.xlabel('number of Folds')
plt.ylabel('mean classification score')
plt.savefig('number-of-cv.png')
plt.show()

plt.plot(scores3, label='over')

```

```


plt.plot(dt_test ,label='test')
plt.legend(loc='best')
plt.title('DecisionTreeClassifier')
plt.xlabel('number of Folds')
plt.ylabel('mean classification score')
plt.savefig('number-of-cv.png')
plt.show()


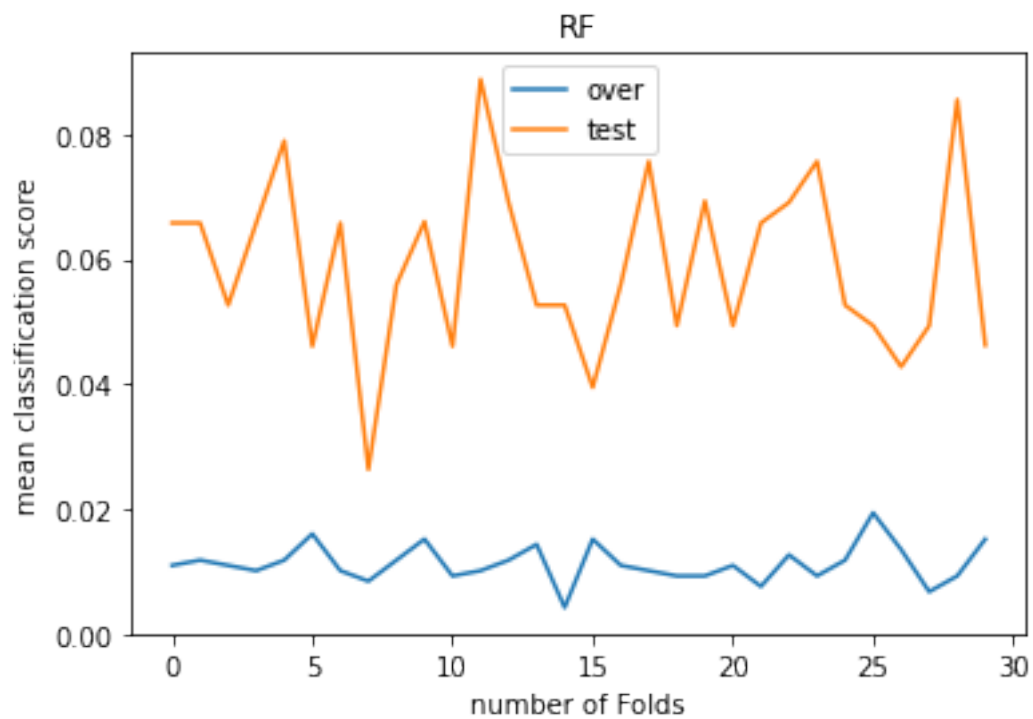
```

Average Brier for Logistic Regression : 0.155 (0.003)

Average Brier Random Forest: 0.011 (0.003)

Average Brier Decision Tree: 0.021 (0.004)





9 Feature Importance

feature importance based on best model Random Forest

```
[ ]: forest = ensemble.RandomForestClassifier(random_state=2002)
      forest.fit(x_over, y_over)
      forest.feature_importances_
```

```
[ ]: array([2.66571716e-02, 1.20702993e-02, 2.20177340e-02, 4.22616346e-02,
          2.93858338e-02, 2.41002413e-02, 2.92718725e-02, 1.07260195e-01,
          3.28002949e-02, 6.42345574e-02, 2.02364333e-01, 1.92517799e-01,
          9.02309290e-02, 6.67051075e-02, 5.72714304e-03, 5.21913643e-03,
          2.43197464e-03, 1.47419096e-03, 2.92736879e-03, 2.72334752e-03,
          1.59999867e-03, 2.55699191e-03, 3.01333341e-03, 1.43272535e-03,
          4.81516105e-03, 4.45981251e-03, 1.75512984e-03, 1.80624571e-03,
          2.64406387e-03, 3.04884972e-03, 2.49432642e-03, 2.46549752e-03,
          1.85752028e-03, 1.47811392e-03, 7.34117665e-04, 1.11980758e-04,
          1.34496734e-03])
```

```
[ ]: # feature importance dataframe
      feat_imp = pd.DataFrame({'Feature': x_over.columns,
                              'Importance': forest.feature_importances_})
      feat_imp_sort = feat_imp.sort_values(by='Importance', ascending=False)
      feat_imp_sort
```

```
[ ]:
```

| | Feature | Importance |
|----|-------------------------------|------------|
| 10 | Total_Trans_Amt | 0.202364 |
| 11 | Total_Trans_Ct | 0.192518 |
| 7 | Total_Revolving_Bal | 0.107260 |
| 12 | Total_Ct_Chng_Q4_Q1 | 0.090231 |
| 13 | Avg_Utilization_Ratio | 0.066705 |
| 9 | Total_Amt_Chng_Q4_Q1 | 0.064235 |
| 3 | Total_Relationship_Count | 0.042262 |
| 8 | Avg_Open_To_Buy | 0.032800 |
| 4 | Months_Inactive_12_mon | 0.029386 |
| 6 | Credit_Limit | 0.029272 |
| 0 | Customer_Age | 0.026657 |
| 5 | Contacts_Count_12_mon | 0.024100 |
| 2 | Months_on_book | 0.022018 |
| 1 | Dependent_count | 0.012070 |
| 14 | Gender_F | 0.005727 |
| 15 | Gender_M | 0.005219 |
| 24 | Marital_Status_Married | 0.004815 |
| 25 | Marital_Status_Single | 0.004460 |
| 29 | Income_Category_\$60K - \$80K | 0.003049 |
| 22 | Education_Level_Unknown | 0.003013 |
| 18 | Education_Level_Graduate | 0.002927 |

| | | |
|----|---------------------------------|----------|
| 19 | Education_Level_High School | 0.002723 |
| 28 | Income_Category_\$40K - \$60K | 0.002644 |
| 21 | Education_Level_Uneducated | 0.002557 |
| 30 | Income_Category_\$80K - \$120K | 0.002494 |
| 31 | Income_Category_Less than \$40K | 0.002465 |
| 16 | Education_Level_College | 0.002432 |
| 32 | Income_Category_Unknown | 0.001858 |
| 27 | Income_Category_\$120K + | 0.001806 |
| 26 | Marital_Status_Unknown | 0.001755 |
| 20 | Education_Level_Post-Graduate | 0.001600 |
| 33 | Card_Category_Blue | 0.001478 |
| 17 | Education_Level_Doctorate | 0.001474 |
| 23 | Marital_Status_Divorced | 0.001433 |
| 36 | Card_Category_Silver | 0.001345 |
| 34 | Card_Category_Gold | 0.000734 |
| 35 | Card_Category_Platinum | 0.000112 |

```
[ ]: from sklearn.feature_selection import SelectFromModel

sel = SelectFromModel(ensemble.RandomForestClassifier(random_state=2002))

#removed if the corresponding importance of the feature values are below the
↪provided threshold parameter
sel.fit(x_over, y_over)
sel.get_support()

selected_feat= x_over.columns[(sel.get_support())]
print(len(selected_feat))
print(selected_feat)
```

```
10
Index(['Total_Relationship_Count', 'Months_Inactive_12_mon', 'Credit_Limit',
      'Total_Revolving_Bal', 'Avg_Open_To_Buy', 'Total_Amt_Chng_Q4_Q1',
      'Total_Trans_Amt', 'Total_Trans_Ct', 'Total_Ct_Chng_Q4_Q1',
      'Avg_Utilization_Ratio'],
      dtype='object')
```

```
[ ]: !sudo apt-get install texlive-xetex texlive-fonts-recommended
      ↪texlive-plain-generic
```

```
Reading package lists... Done
Building dependency tree
Reading state information... Done
texlive-fonts-recommended is already the newest version (2017.20180305-1).
texlive-plain-generic is already the newest version (2017.20180305-2).
texlive-xetex is already the newest version (2017.20180305-1).
The following package was automatically installed and is no longer required:
```

```
libnvidia-common-460
Use 'sudo apt autoremove' to remove it.
0 upgraded, 0 newly installed, 0 to remove and 5 not upgraded.
```

```
[64]: !jupyter nbconvert --to pdf /Customers_Attraction.ipynb
```

```
[NbConvertApp] WARNING | pattern '/content//Customers_Attraction.ipynb' matched
no files
This application is used to convert notebook files (*.ipynb)
to various other formats.
```

```
WARNING: THE COMMANDLINE INTERFACE MAY CHANGE IN FUTURE RELEASES.
```

Options

=====

The options below are convenience aliases to configurable class-options,
as listed in the "Equivalent to" description-line of the aliases.

To see all configurable class-options for some <cmd>, use:

<cmd> --help-all

--debug

set log level to logging.DEBUG (maximize logging output)

Equivalent to: [--Application.log_level=10]

--show-config

Show the application's configuration (human-readable format)

Equivalent to: [--Application.show_config=True]

--show-config-json

Show the application's configuration (json format)

Equivalent to: [--Application.show_config_json=True]

--generate-config

generate default config file

Equivalent to: [--JupyterApp.generate_config=True]

-y

Answer yes to any questions instead of prompting.

Equivalent to: [--JupyterApp.answer_yes=True]

--execute

Execute the notebook prior to export.

Equivalent to: [--ExecutePreprocessor.enabled=True]

--allow-errors

Continue notebook execution even if one of the cells throws an error and
include the error message in the cell output (the default behaviour is to abort
conversion). This flag is only relevant if '--execute' was specified, too.

Equivalent to: [--ExecutePreprocessor.allow_errors=True]

--stdin

read a single notebook file from stdin. Write the resulting notebook with
default basename 'notebook.*'

Equivalent to: [--NbConvertApp.from_stdin=True]

--stdout

Write notebook output to stdout instead of files.
 Equivalent to: [--NbConvertApp.writer_class=StdoutWriter]

--inplace
 Run nbconvert in place, overwriting the existing notebook (only relevant when converting to notebook format)
 Equivalent to: [--NbConvertApp.use_output_suffix=False
 --NbConvertApp.export_format=notebook --FilesWriter.build_directory=]

--clear-output
 Clear output of current file and save in place, overwriting the existing notebook.
 Equivalent to: [--NbConvertApp.use_output_suffix=False
 --NbConvertApp.export_format=notebook --FilesWriter.build_directory=
 --ClearOutputPreprocessor.enabled=True]

--no-prompt
 Exclude input and output prompts from converted document.
 Equivalent to: [--TemplateExporter.exclude_input_prompt=True
 --TemplateExporter.exclude_output_prompt=True]

--no-input
 Exclude input cells and output prompts from converted document.
 This mode is ideal for generating code-free reports.
 Equivalent to: [--TemplateExporter.exclude_output_prompt=True
 --TemplateExporter.exclude_input=True]

--log-level=<Enum>
 Set the log level by value or name.
 Choices: any of [0, 10, 20, 30, 40, 50, 'DEBUG', 'INFO', 'WARN', 'ERROR', 'CRITICAL']
 Default: 30
 Equivalent to: [--Application.log_level]

--config=<Unicode>
 Full path of a config file.
 Default: ''
 Equivalent to: [--JupyterApp.config_file]

--to=<Unicode>
 The export format to be used, either one of the built-in formats ['asciidoc', 'custom', 'html', 'latex', 'markdown', 'notebook', 'pdf', 'python', 'rst', 'script', 'slides'] or a dotted object name that represents the import path for an `Exporter` class
 Default: 'html'
 Equivalent to: [--NbConvertApp.export_format]

--template=<Unicode>
 Name of the template file to use
 Default: ''
 Equivalent to: [--TemplateExporter.template_file]

--writer=<DottedObjectName>
 Writer class used to write the results of the conversion
 Default: 'FilesWriter'

which will convert mynotebook.ipynb to the default format (probably HTML).

You can specify the export format with `--to``.

Options include ['asciidoc', 'custom', 'html', 'latex', 'markdown', 'notebook', 'pdf', 'python', 'rst', 'script', 'slides'].

```
> jupyter nbconvert --to latex mynotebook.ipynb
```

Both HTML and LaTeX support multiple output templates. LaTeX includes

'base', 'article' and 'report'. HTML includes 'basic' and 'full'.

You

can specify the flavor of the format used.

```
> jupyter nbconvert --to html --template basic mynotebook.ipynb
```

You can also pipe the output to stdout, rather than a file

```
> jupyter nbconvert mynotebook.ipynb --stdout
```

PDF is generated via latex

```
> jupyter nbconvert mynotebook.ipynb --to pdf
```

You can get (and serve) a Reveal.js-powered slideshow

```
> jupyter nbconvert myslides.ipynb --to slides --post serve
```

Multiple notebooks can be given at the command line in a couple of different ways:

```
> jupyter nbconvert notebook*.ipynb
```

```
> jupyter nbconvert notebook1.ipynb notebook2.ipynb
```

or you can specify the notebooks list in a config file, containing::

```
c.NbConvertApp.notebooks = ["my_notebook.ipynb"]
```

```
> jupyter nbconvert --config mycfg.py
```

To see all available configurables, use `--help-all``.